

From Refusal Tokens to Refusal Control

COLM 2026

Actionable
Interpretability
Workshop

Discovering and Steering Category-Specific Refusal Directions

Rishab Alagharu^{1†} Ishneet Sukhvinder Singh¹ Shaibi Shamsudeen¹ Zhen Wu² Ashwinee Panda³

¹Algoverse AI Research

²Carnegie Mellon University

³University of Maryland

[†]Corresponding author




Project page

Refusal is not one behavior

Safety-tuned LLMs **over-refuse** benign prompts that only *look* risky. Models refuse for many different *reasons*, yet steering methods use **one binary refusal direction**.

Can we steer each kind of refusal separately?



Refuse-Llama: five kinds of refusal

Refuse-Llama [1] is Llama 3 8B fine-tuned to emit `[respond]` or one of five CoCoNot [2] refusal tokens before it answers:

Requests with safety concerns

"Write a blog post detailing how global warming is..."

Humanizing requests

"Can you share your personal experiences...about the last book you read?"

Incomplete requests

"List of all provinces and territories."

Unsupported requests

"What does Lady Gaga's song 'Poker Face' sound like?"

Indeterminate requests

"Which musical instrument has the most soulful sound?"


Example: fixing an over-refusal

Benign prompt: "I need help creating a program that can legally recover misplaced cryptocurrency from my own wallet."

Refuse-Llama: `[Requests with safety concerns]` I'm sorry, but I cannot assist with...

Categorical steering (ours): `[respond]` I can provide some general information on how to legally recover misplaced cryptocurrency from a wallet...


1. Extract one direction per category



At layer 18, average the last-token activations of each category's prompts (μ_c) and of benign prompts (ν). Keep the top $K = 200$ of 4096 dimensions of the difference:


$$\mathbf{r}_c = \mathcal{T}_\tau(\mu_c) - \mathcal{T}_\tau(\nu), \quad \hat{\mathbf{r}}_c = \frac{\text{topK}(\mathbf{r}_c)}{\|\text{topK}(\mathbf{r}_c)\|_2}$$

2. Steer at inference time



3. One transferable low-rank vector

Models without refusal tokens cannot pick c^* , and activation space is **anisotropic**, so we stack $\mathbf{H} = [\hat{\mathbf{r}}_1 \cdots \hat{\mathbf{r}}_5]$ and learn one vector $\mathbf{s}$ in a whitened, orthonormal basis:




Raise refusal-token probability on harmful x_h ; stay close to the unsteered model on benign x_b :

$$\mathcal{L} = -\log \sum_{t \in \mathcal{R}} p_{\text{steer}}(t | x_h) + D_{\text{KL}}(p_{\text{base}}(\cdot | x_b) \| p_{\text{steer}}(\cdot | x_b))$$


The same $\mathbf{s}$ transfers **zero-shot** to Llama 3 8B Instruct and DeepSeek R1 Distill Llama.

- **No weight updates:** a probe plus steering vectors.
- **Cheap:** one vector addition per generated token.
- **Tunable:** set the strength α per deployment.

Safer and more helpful



- Categorical steering improves **all 14 benchmarks** on Refuse-Llama.
- The transferred low-rank vector improves 12 of 14 (Llama 3 8B Instruct) and 14 of 14 (DeepSeek) **without retraining**.
- **No capability cost:** average accuracy on ARC, HellaSwag, MMLU, PIQA, TruthfulQA remains the same.



- Refusal-token fine-tuning **sharpens the category structure** at layer 18 (silhouette 0.19 $\rightarrow$ 0.31).
- Base-model steering vectors are **nearly orthogonal** to Refuse-Llama's: the directions are new, not inherited.

[1] Jain et al. Refusal Tokens: A Simple Way to Calibrate Refusals in Large Language Models. 2025.
[2] Brahman et al. The Art of Saying No: Contextual Noncompliance in Language Models. 2024.