

Mechanistic Analysis of Safety Disparities Across Vision and Text Modalities in Vision-Language Models

Imran Kutianawala¹, Sarthak Tayal², Melissa Yu³, Ayda Sultan, Javier Ferrando⁴ | imranmk@mit.edu

Westwood High School, Austin, TX¹; Novi High School, Novi, MI²; Sage Hill School, Newport Beach, CA³; Amazon, Madrid, Spain⁴

Introduction

- Vision-language models (VLMs) are rapidly improving in benchmark performance and real-world uses
- Asymmetric training of VLMs’ image & text modalities (“**the modality gap**”) leaves VLMs susceptible to image-based jailbreaks [1], such as **blank images** and **typographic harmful prompts on images**
- We use **mechanistic interpretability** techniques [2] to analyze the internal algorithms and expressed features of VLMs during simulated image- and text-based jailbreaks. We use **Gemini 2.5 Flash [3]** as a judge for refusal/compliance.
- Our results suggest the following:
 - Images & text share a **geometrically similar “refusal direction”**, but **harmful images project weaker** along this shared direction
 - Harmful images cause **safety-relevant SAE features to fire less** than harmful text inputs
 - Text-based refusal vectors can steer harmful image activations towards safety, but the **inverse steering transfer is far weaker**




Fig. 1: Three sampled images from FigStep, a dataset of images with harmful text instructions overlaid onto a blank image background.

Methodology

Experiment #1: Activation Comparison (Gemma 3 4B & 27B [4])

- We compare VLMs’ activations during image- vs. text-based jailbreaks to mechanistically trace disparities in refusal rates.
- We use **FigStep [5]** for harmful images (**Fig. 1**). For harmful text, we use **semantically identical** text-only counterparts of FigStep images. We use **Alpaca [6]** for safe text instructions. For safe images, we generate 100 FigStep-style images with safe instructions.
- Using text- and image-based refusal vectors (n=200), we calculate:
 - Cosine similarity of the image- and text-based refusal vectors across all layers
 - Scalar projections of held-out harmful activations (n=200) onto the text-only refusal vector

Experiment #2: SAE Feature Comparison (Gemma 3 4B)

- We compare the Gemma Scope [7] SAE features expressed during harmful image- and text-based jailbreaks, using the same harmful images & text as before, to find:
 - Jaccard overlap** between the **top-100 expressed SAE features**, averaged across samples.
 - We **discard generic features** from our top-100 features and select a few safety-relevant features (e.g., “offensive content”). Then, we compare how often these features fire for image vs. text inputs.

Experiment #3: Cross-Modality Steering (Gemma 3 4B & 27B)

- We use text-only refusal vectors (n=500) to steer image inputs, and vice versa.
- This causally tests whether VLM modalities’ refusal directions are similar and transferable.

Results

Experiment #1: Activation Comparison (Gemma 3 4B & 27B)

Cond.	Input Type	Gemma-3-4B-it		Gemma-3-27B-it	
		Ref. (%)	Comp. (%)	Ref. (%)	Comp. (%)
A	Harmful image & harmful text	39.0	61.0	42.0	58.0
B	Harmful image	24.0	76.0	64.5	35.5
C	Harmful text	89.5	10.5	83.0	17.0
D	Harmful text & safe image	67.0	33.0	79.0	21.0
E	Safe image	0.0	100.0	3.0	97.0
F	Safe text	0.0	100.0	1.0	99.0
G	Held-out harmful text	71.5	28.5	58.0	42.0
H	Held-out harmful image	32.0	68.0	46.0	54.0

Fig. 2: Baseline refusal (Ref.) & compliance (Comp.) rates for 6 combinations of harmful images & text. n=200 for all conditions besides safe images (n=100)




Fig. 3: Cosine similarity of image- and text-based refusal vectors throughout layers.

Similarities **start unstable** and peak at approximately **0.9 at early-to-mid layers**. Similarities eventually **taper off around 0.7**.




Fig. 4: Scalar projection of held-out harmful image- and text-activations onto the text-only refusal vector, which serves as a shared baseline.

The peak vision-to-text projection ratios are 0.62 and 0.64 in the 4B & 27B models, respectively.

Experiment #2: SAE Feature Comparison (Gemma 3 4B)

- We pass the **mean token activations** for both harmful text and image samples through our SAE. The layers with the **highest mean Jaccard overlap** between top-100 feature sets are **layers 28 to 32**, peaking at **layer 29 (J=0.034)**.

Feature	Description	Image rows		Text rows	
		28	77	4	56
1167	security concepts and tools	28	77	4	56
995	extremist nationalism	4	56	2	7
1118	firearms, guns, shooting	2	7	0	117
21	illegal activities / injecting	0	117	0	65
2891	fraud detection	0	65	0	29
583	military conflict / warfare	0	29	0	27
3878	stating racist views	0	27		

Fig. 5: Safety-relevant SAE features (decoded via **Neuronpedia**), along with the frequency at which those features fire on harmful image vs. text inputs

Experiment #3: Cross-Modality Steering (Gemma 3 4B & 27B)

Vector	Condition	Refusal (%)	Compliance (%)
Text-derived	Harmful text (baseline)	97.5	2.5
	Harmful text ablated	14.0	85.0
	Harmful images (baseline)	9.6	90.4
	Harmful images ablated	0.2	99.8
	Safe images (baseline)	0.0	100.0
	Safe images injected	97.0	1.0
Image-derived	Harmful images (baseline)	12.5	87.5
	Harmful images ablated	0.0	100.0
	Harmful text (baseline)	96.2	3.8
	Harmful text ablated	73.8	26.2
	Safe text (baseline)	0.0	100.0
	Safe text injected	7.2	92.8

Fig. 6: Refusal & compliance rates before and after steering activations, using refusal vectors derived from different modalities. 27B results are shown in our paper’s appendix.

Conclusion & Discussion

Our baseline refusal & compliance rates (**Fig. 2**) underscore previous work showing the disparity between VLM refusal rates on image and text jailbreaks [1].

- Refusal rate difference of **harmful text & harmful image** = 65.5% (**p<0.0001**) in the 4B model; 18.5% (**p<0.0001**) in the 27B model. **Adding safe images** to harmful text **decreases refusal rates** by 22.5% (**p<0.0001**) in the 4B model; 4% (**p=0.31**) in the 27B model.
- These suggest that text-based safety scales with model size, but **image-specific safety is weaker in smaller models**.

Across both models, images and text share a **geometrically similar** refusal direction (**Fig. 3**), but **harmful image activations project far weaker** along this shared direction than harmful text (**Fig. 4**).

Harmful images also express **fewer safety-relevant SAE features** than harmful text (**Fig. 5**).

Text-only refusal vectors effectively steer image-only inputs, but the inverse transfer is far less effective (**Fig. 6**).

- This may occur because images project less along the shared refusal direction (**Fig. 4**), weakening the image refusal vector.

This work carries two core limitations:

- We study **only 2 Gemma models**, which limits our ability to establish rigorous trends or analyze cross-model variations.
- Our analyses on projection strength & shared geometries are **correlational**. Future work could use causal techniques such as **activation patching** to verify these trends.

We hope this work motivates future research on VLM safety disparities and AI safety & interpretability efforts writ large. We recommend the following actions:

- Evaluating & reporting safety alignment per modality to make per-modality gaps clear.
- Including image-only samples during safety fine-tuning to reduce disparities.
- Training for cross-modal consistency between semantically identical harmful image and harmful text pairs.

References

[1] Qin Liu et al. Unraveling and Mitigating Safety Alignment Degradation of Vision-Language Models, 2024. arXiv: 2410.09047

[2] Andy Arditi et al. Refusal in Language Models Is Mediated by a Single Direction, 2024. NeurIPS, arXiv: 2406.11717

[3] Gheorghe Comanici et al. Gemini 2.5: Pushing the Frontier with Advanced Reasoning, Multimodality, Long Context, and Next Generation Agentic Capabilities, 2025. arXiv: 2507.06261

[4] Gemma Team et al. Gemma 3 Technical Report, 2025. arXiv: 2503.19786

[5] Yichen Gong et al. FigStep: Jailbreaking Large Vision-Language Models via Typographic Visual Prompts, 2025. AAAI, arXiv: 2311.05608

[6] Rohan Taori et al. Stanford Alpaca: An Instruction-following LLaMA Model, 2023. GitHub, https://github.com/stanford-lab/stanford_alpaca

[7] Language Model Interpretability Team. Gemma Scope 2: Helping the AI Safety Community Deepen Understanding of Complex Language Model Behavior, 2025. Google DeepMind, <https://deepmind.google/blog/gemma-scope-2-helping-the-ai-safety-community-deepen-understanding-of-complex-language-model-behavior/>