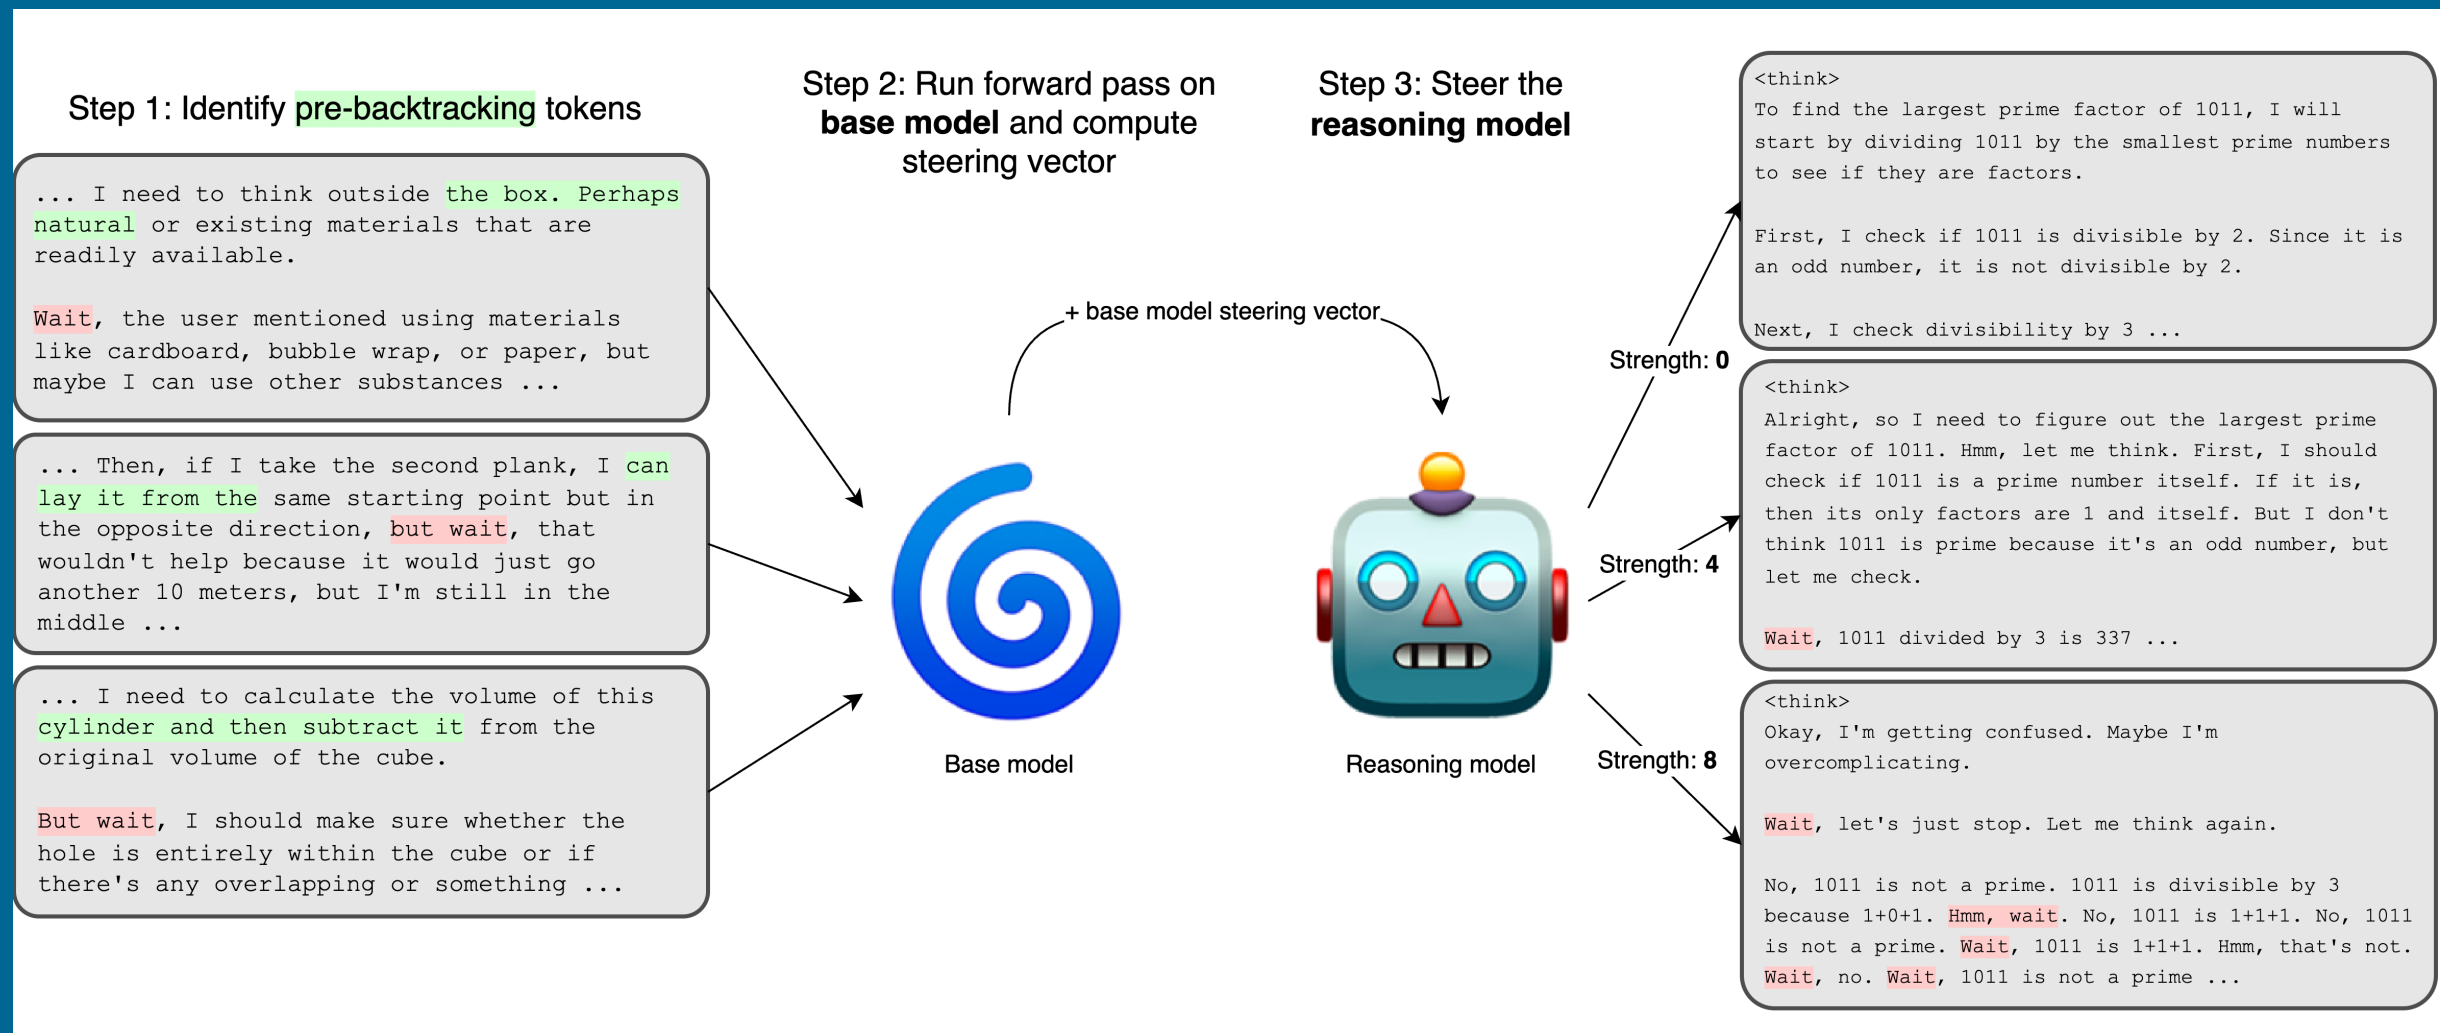


We Found Directions in *Base Models* That Trigger Reasoning Behavior in *Thinking Models*.



Reasoning-Finetuning Repurposes Latent Representations in Base Models

Jake Ward*, Chuqiao Lin[†], Neel Nanda *Independent [†]Rudolf Peierls Centre for Theoretical Physics, Oxford, UK, OX1 3PU

We study the mechanism of backtracking from reasoning FT

Backtracking: An emergent behavior in *reasoning* models where the model explores alternative strategies after progressing down a reasoning path or proposing a candidate answer. It accounts for a substantial fraction of the accuracy gap between base and reasoning-fine-tuned models.

Key questions:

- How does backtracking emerge during reasoning fine-tuning?
- Are these capabilities learned from scratch or built on existing representations?

Our method: training steering vectors for backtracking

Training Setup: We study backtracking behavior with **Llama-3.1-8B** (base model) and **DeepSeek-R1-Distill-Llama-8B** (reasoning fine-tuned model). We query LLM judges to identify backtracking events in 300 reasoning traces. The steering vectors are derived using the **Difference-of-Means (DoM)** method. Additionally, we

- Derive steering vectors with a **negative token position offset** - The directions extracted from some tokens before actual backtracking suggests that they are causally relevant to the model's decision to backtrack.
- Use **base model activations**: The steering vectors are derived from the residue stream in the forward pass of reasoning traces on both the *base* model and the *finetuned* model.

Eval Setup: The derived steering vectors are evaluated by their ability to induce future backtracking events - we measure the frequency of backtracking tokens ("wait", "but", "Hmm", etc) in rollouts of text generation after steering.

Main Results

Optimal Steering Parameters (Fig. 1)

- Optimal token offset:** -13 to -8 tokens **before** backtracking event - typically covers the beginning of the sentence prior to backtracking.
- Optimal steering layer:** Most effective around layer 10, consistent with previous results (Venhoff, et. al. 2025)

Base-Derived Vectors Induce Backtracking When Applied To The Reasoning Model (Main Fig, Fig 2.)

- High cosine similarity (0.74)** between base-derived and reasoning-derived steering vectors - suggesting **shared representations** between base and reasoning models

- Base model never backtracks** when steered with base/reasoning-derived vectors




Fig 1. The effect of steering as a function of token window offset and steering vector magnitude. (Layer-10 residue stream of the reasoning model)

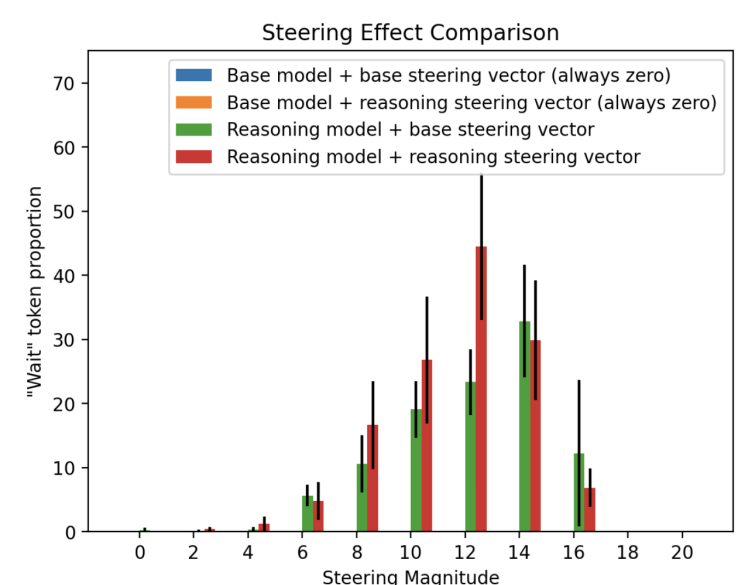



Fig 2. Proportion of backtracking-related tokens generated by both base and reasoning models when steered with base-derived or reasoning-derived steering vectors.

Interpreting the steering directions (is hard!)

Logit lens analysis: Both base-derived and reasoning derived vectors **do not directly boost** backtracking token probabilities - they cannot be explained by token-level attributes, suggesting they capture **more abstract concepts causally relevant for backtracking**



Logit lens experiments on steering vectors trained on the base/fine-tuned model.